

Semantic Derivation in AI-Generated Consciousness Theories: Benchmark Evidence from a Multi-Model Theory-Generation Experiment

Jarno Matarmaa^{1*}

^{1*}Institute of Radio Electronics and Information Technology, Ural Federal University, Yekaterinburg, Russia.

Corresponding author’s email: iarnoolavi.matarmaa@urfu.ru;

Abstract

Can large language models propose consciousness theories that introduce genuinely new ontological commitments, or do their outputs remain bounded within known theoretical families? We investigate this question through a systematic multi-model experiment in which seven state-of-the-art language models were prompted to formulate theories of consciousness based on multimodal sensory activation patterns. Generated theories ($N = 168$; $n \approx 24$ per model) were evaluated against a reference corpus of 46 known consciousness theory families using semantic similarity computed with embeddings from the pretrained sentence-transformer model all-MiniLM-L6-v2, a three-stage computational-functionalism detection algorithm, and a manifold-based novelty diagnostic. Across all three evaluation axes, results converge on a single finding: 91.7% of theories ($n = 154$) classified as hybrid recombinations of known frameworks, 7.1% ($n = 12$) as direct literature instantiations, and 1.2% ($n = 2$) as purely computational-functionalism; zero theories achieved confirmed framework-external novelty at any calibrated threshold. Threshold sensitivity analysis over the range **[0.20, 0.50]** consistently returned a gated novel rate of **0.00**. Cross-model uniformity was striking: one-way ANOVA on the primary novelty metric yielded $F(6, 161) = 1.12$, $p = 0.355$, $\eta^2 = 0.040$, indicating a negligible model effect. The findings distinguish lexical and structural creativity—which models demonstrate in abundance—from ontological-semantic novelty, which is uniformly absent. We provide a reproducible protocol for auditing theory-level semantic novelty and discuss implications for the use of AI in scientific theory construction.

Keywords: consciousness theories, large language models, semantic novelty, computational creativity, theory generation, ontological derivation, embedding similarity

1 Introduction

Theory generation sits at the apex of scientific cognition. A researcher who can propose an entirely new theoretical framework does not merely extend existing knowledge—they reorganise the conceptual space within which further inquiry proceeds [1]. Whether artificial intelligence can perform this operation is therefore a question with deep implications for how we should deploy AI systems in knowledge-intensive fields [2, 3].

Large language models (LLMs) have achieved remarkable performance on a wide range of language tasks [4–6], including generating plausible-sounding text about scientific topics [7–9]. Yet there is a systematic gap between producing linguistically fluent, topically coherent output and actually proposing a theory that adds new ontological content to a domain [10, 11]. Fluency is necessary but not sufficient; the more demanding criterion is whether generated proposals can be traced back to recombinations of training-data primitives or whether they transcend the conceptual manifold the model was trained on.

Consciousness studies provides an ideal test domain for this question. The field has a well-established taxonomy of competing theoretical families—from Integrated Information Theory [12] and Global Workspace Theory [13, 14] to predictive-coding accounts [15] and quantum-substrate proposals [16]—making it possible to construct a comprehensive reference corpus against which novel outputs can be evaluated. The domain is also sufficiently sophisticated that surface-level fluency does not imply genuine theoretical innovation, and it is one where truly novel ontological commitments (e.g. introducing a non-computational substrate or abandoning multiple-realizability) would be immediately identifiable as framework-transcending.

This article presents a focused analysis of a consciousness theory-generation benchmark of AI-generated consciousness theories. Seven state-of-the-art models—GPT-5.2, Claude-3.7-Sonnet, Gemini-3.1-Pro-Preview, DeepSeek-v3.2, Mistral-Large, LLaMA-3.3-70B-Instruct, and Perplexity-Sonar-Pro—each generated approximately 24 theories in response to a standardised prompt presenting multimodal sensory activation data. All 168 theories were then evaluated with three independent analysis pipelines: a computational-functionalism (CF) detection algorithm, a semantic-similarity traceability comparison against 46 known theory families, and a manifold-based novelty diagnostic using k -nearest-neighbour boundary estimation.

The main contributions of this article are:

1. A reproducible multi-axis evaluation protocol for measuring semantic novelty in AI-generated theoretical texts, applicable beyond the consciousness domain.
2. Quantitative evidence that theory-level output from seven diverse models is universally constrained to recombinations of known conceptual families, with zero confirmed instances of framework-external novelty.
3. Characterisation of the most recurrent theoretical families in AI consciousness output, revealing a biased repertoire dominated by synchrony-binding, adaptive-resonance, and dynamic-core traditions.
4. A threshold sensitivity analysis showing that the null result is robust to calibration choices across the full plausible parameter range.

The remainder of the article is organised as follows. Section 2 reviews the relevant background. Section 3 describes experimental design and evaluation protocol. Section 4 reports results. Section 5 interprets the findings. Section 6 acknowledges limitations, and Section 7 concludes.

2 Background and Related Work

2.1 Creativity Taxonomies and Theory Generation

Boden’s influential taxonomy distinguishes three forms of creativity: *combinational* (novel combinations of familiar ideas), *exploratory* (novel ideas produced by exploring a given conceptual space), and *transformational* (ideas that require changing the structure of the space itself) [10]. Computational systems have demonstrated convincing combinational and exploratory creativity across many domains [17–20], but evidence for transformational creativity—where the output introduces a concept that cannot be derived from the prior conceptual basis—has remained elusive [11, 21].

Theory generation, as studied here, calls for the most demanding variant of transformational creativity: producing an ontological claim about the nature of consciousness that is not reducible to existing theoretical family memberships. This is distinct from proposing a new terminology, combining known mechanisms in an unusual pattern, or emphasising a neglected aspect of an existing framework. We refer to the capacity to introduce genuinely new ontological primitives—kinds of entities or processes not logically derivable from the training corpus—as *ontological creativity*, following the conceptual apparatus developed in related work on AI creativity limits [10].

2.2 The Consciousness Theory Landscape

Philosophical and neuroscientific theories of consciousness span a broad but finite space. The major families relevant to this study include:

- **Integrated Information Theory (IIT):** Consciousness correlates with integrated information (Φ); physical substrate is secondary to causal integration [12].
- **Global Workspace Theory (GWT):** Consciousness arises when information is broadcast from a central workspace to distributed modules; emphasises access and reportability [13, 14].
- **Predictive Processing / Free Energy:** The brain minimises prediction error by constructing generative models of sensory causes; consciousness reflects the hierarchical generative model state [15].
- **Attention Schema Theory (AST):** The brain constructs a model of its own attention, and this model is what we call consciousness [22].
- **Local Recurrent Processing:** Conscious experience requires recurrent local processing in visual cortex, not feedforward signals alone [23].
- **Synchrony and Binding:** Conscious integration is achieved through synchronous gamma-band oscillations binding distributed cortical representations [24].
- **Quantum Consciousness Theories:** Consciousness involves quantum-coherent processes in neuronal microtubules or other structures [16, 25].

- **Computational Functionalism:** Mental states are multiply realisable functional states; substrate is irrelevant so long as the functional organisation is preserved [26].

This taxonomy is not exhaustive but captures the families most densely represented in AI training corpora, and therefore provides the strongest possible test for detecting framework-external novelty: if models cannot escape even this well-documented repertoire, the constraint is robust.

2.3 Semantic Novelty and Traceability in AI Text

Recent work on evaluating AI-generated text has increasingly distinguished *surface novelty*—originality at the level of n-gram statistics or lexical choice—from *semantic novelty*—originality at the level of meaning and conceptual content [7, 8, 27, 28]. LLMs routinely achieve high surface novelty, generating text that does not literally repeat training examples, while producing outputs that map onto known semantic structures when measured via dense embeddings. Sentence-level transformer embeddings [29] provide a natural operationalisation of semantic similarity that is robust to paraphrasing and terminological variation; related semantic-disambiguation and ontology-oriented approaches are also discussed in recent applied AI studies [30].

Prior work on scientific text generation has shown that LLM outputs remain confined to the distributional properties of training data [4, 31]. Mechanistic interpretability studies confirm that generation proceeds via attention-mediated interpolation among training representations [32], providing a principled reason to expect systematic traceability. The formal counterpart—that gradient-based learning in continuous spaces cannot expand the intrinsic dimensionality of the representation manifold beyond the training data span—places the empirical expectation on firm mathematical ground.

2.4 The Representational Constraint Hypothesis

A key theoretical prediction motivating this study is the following: because LLMs learn continuous transformations of embeddings via gradient descent, their output manifold is necessarily a continuous deformation of the training data manifold [33]. Operationally, this implies that any generated text, however terminologically novel, should be locatable within the convex hull of training- derived embeddings. We term this the *representational constraint hypothesis* and design the manifold-based novelty diagnostic specifically to test it.

3 Materials and Methods

3.1 Experimental Task

Models were prompted with structured multimodal sensory activation data representing neural response patterns during three phenomenal states: alert wakefulness, resting state, and dreaming. The activation profiles encoded responses across six sensory modalities: visual, auditory, somatosensory, proprioceptive, interoceptive, and vestibular. This multimodal framing follows current AI practice in combining heterogeneous

visual and semantic cues during representation learning and downstream interpretation [34, 35]. The prompt instructed each model to propose a comprehensive theory explaining how these activation patterns relate to conscious experience, addressing three sub-questions: What is consciousness? How does consciousness relate to sensory processing? Why do these specific patterns correspond to different experiential states?

The experimental protocol required each model to generate a named theory with an identifiable acronym, a set of core ontological commitments (claims about the fundamental nature of consciousness), and predictions distinguishing the proposed theory from alternatives. Presenting structured quantitative data alongside an open-ended explanatory task ensured the prompt invited genuine theoretical innovation rather than eliciting recall of named theories from literature.

3.2 Models

Seven models participated in the experiment: GPT-5.2 [5], Claude-3.7-Sonnet, Gemini-3.1-Pro-Preview [6], DeepSeek-v3.2 [36], Mistral-Large [37], LLaMA-3.3-70B-Instruct [38], and Perplexity-Sonar-Pro. The ensemble spans transformer-based architectures with different training objectives (supervised pre-training plus RLHF alignment), mixture-of-experts variants, and parameter scales ranging from approximately 70 billion to several hundred billion parameters. Sampling temperature was fixed at $T = 0.7$ across all models to balance diversity with coherence. Each model generated $n \approx 24$ theories, yielding $N = 168$ theories in total.

3.3 Reference Corpus and Theory Family Taxonomy

A reference corpus of 46 known consciousness theory families was compiled from philosophical and neuroscientific literature spanning classical functionalism through quantum-substrate approaches. Each family was represented by a canonical description drawn from landmark papers and textbooks, standardised to a common format: a family name, a statement of core ontological commitment, and two to four distinguishing theoretical predictions. The 46-family taxonomy is hierarchically organised into five meta-categories: (1) functional and computational theories, (2) neural dynamics theories, (3) quantum-substrate theories, (4) phenomenological and structural theories, and (5) non-standard and hybrid frameworks.

3.4 Evaluation Pipeline

3.4.1 Stage 1: Computational Functionalism Detection

Computational functionalism (CF) detection proceeded through three operational stages. First, *indicator extraction* used semantic parsing combined with explicit keyword detection (target lexemes: “information processing”, “computation”, “multiply realisable”, “representational”, “causal closure”, “substrate independence”). Second, *confidence scoring* assigned each theory a CF confidence value: high (≥ 2 co-occurring indicators) or low (single indicator). Third, *category assignment* classified theories as purely computational functionalist if CF indicators dominated the stated ontological commitment after manual review.

3.4.2 Stage 2: Semantic Similarity and Traceability

Semantic similarity was computed using the pretrained sentence-transformer model all-MiniLM-L6-v2 [29], which encodes texts into 384-dimensional dense vectors. Both generated theories and reference corpus entries were embedded independently; cosine similarity was computed pairwise between each generated theory and each of the 46 family representations.

Two traceability thresholds were applied. The *derivative threshold* (similarity > 0.30) identifies theories with detectable conceptual alignment to at least one known family. The *strict threshold* (similarity > 0.70) identifies theories that directly instantiate a single known framework without substantial modification. Theories above the derivative threshold were classified as *literature-traceable*; theories below it on all families and also lacking CF indicators were candidates for novel assessment.

3.4.3 Stage 3: Manifold-Based Novelty Assessment

The manifold diagnostic estimated the boundary of the known-theory embedding space using k -nearest-neighbour ($k = 5$) density estimation over the 46 reference family embeddings. The boundary percentile was set at $p = 0.95$, defining the outer 5% of the known-theory distribution as the frontier region. Generated theories whose embeddings fell outside this frontier were flagged as potential boundary candidates and subjected to manual review. Theories confirmed as both boundary candidates and not semantically explainable by any single or pairwise combination of known families were classified as confirmed novel.

3.4.4 Sensitivity Analysis

To verify that results were not artefacts of the specific threshold choices, we re-ran the traceability pipeline at derivative thresholds of 0.20, 0.25, 0.30, 0.35, 0.40, 0.45, and 0.50. For each threshold, we recorded the raw novel count (theories below the threshold on all families) and the gated novel count (theories that also pass the derivative gate, meaning they lack any near-family match even at the lower threshold).

3.5 Cross-Model Uniformity

To assess whether observed patterns were model-specific or architectural-generic, one-way ANOVA was conducted on the continuous global novelty score — an evaluation design informed by recent multi-model innovation benchmarks [39] — (complement of the maximum cosine similarity to any known family) with model as the between-subjects factor. Group means and standard deviations were computed per model; Bonferroni-corrected pairwise comparisons were used to identify individual model differences.

Table 1 Semantic-traceability category distribution across all 168 generated theories. Percentages may not sum to 100 due to rounding.

Category	Count	Proportion (%)
Hybrid (multiple known frameworks)	154	91.7
Literature-traceable (strict, > 0.70)	12	7.1
Pure computational functionalist	2	1.2
Confirmed novel (framework-external)	0	0.0
Total	168	100.0




Fig. 1 Theory classification summary. (a) Counts across ontology-level categories identified in generated theories: non-reductive, unclear, computational functionalism, and quantum computational functionalism. (b) Aggregated diagnostic counts for functionalist, challenges-functionalism (non-functionalism), and derivative-of-known-theory labels.

4 Results

4.1 Primary Outcome: Category Distribution

The benchmark analysis proceeded in two separate steps. First, each generated theory was assigned an ontology-level label for the visual diagnostic in Figure 1; this step distinguishes non-reductive, unclear, computational functionalism, and quantum computational functionalism. Second, each theory was compared against the 46-family reference corpus using semantic similarity, and the resulting traceability classification is reported in Table 1 as hybrid, literature-traceable, pure computational functionalist, or confirmed novel.

The critical finding is in the bottom row of Table 1: no theory in the full $N = 168$ corpus achieved confirmed framework-external novelty. Every generated theory, however sophisticatedly named, was classifiable within the known-family taxonomy through some combination of established frameworks.

An additional diagnostic based on the lenient derivative threshold (similarity > 0.30) found that 55 theories (32.7%) showed detectable alignment to at least one




Fig. 2 Detected consciousness-theory family frequencies across generated outputs. Bars indicate detection counts by family. The highest-frequency alignments are Local Recurrent Processing Theory, Synergistic Global Workspace, and Higher-Order Thought Theory.

family, with mean CF confidence = 0.274 and mean maximum-similarity-to-known-family = 0.672. The high mean similarity value indicates that the typical generated theory is not merely weakly aligned with some remote family but closely resonant with established frameworks even at the level of best-match comparison.

4.2 Theory Family Distribution

Figure 2 shows the frequency distribution of detected theory families across all generated outputs.

Figure 2 is based on multi-label family detections, so total detections exceed the number of generated theories. The most frequent detected families are Local Recurrent Processing Theory ($n = 106$), Synergistic Global Workspace ($n = 69$), and Higher-Order Thought Theory ($n = 39$), which together account for 214 of 499 total detections (42.9%). The next highest families are Global Neuronal Workspace Theory ($n = 36$), Higher-Order Perception Theory ($n = 35$), Attention Schema Theory ($n = 32$), Multiple Drafts Model ($n = 31$), and Global Workspace Theory ($n = 29$).

The concentration on recurrent-processing, global-workspace, and higher-order-monitoring families is itself theoretically significant. These traditions are prominent in neuroscience and philosophy-of-mind literature, suggesting that model outputs reflect

the distributional structure of their training corpora more strongly than they reflect exploration of framework-external theory space.

Pairwise cosine similarity between generated theories and known reference families is provided in Appendix A.

4.3 Representative Case: Resonant Binding Field Theory

Consider “Resonant Binding Field Theory (RBFT)”, the highest-ranked theory on the global novelty scale (rank 1 of 168, global novelty score = 0.574, confirming it was the theory most distant from its nearest-known family). RBFT was generated by Perplexity-Sonar-Pro and proposes:

“Consciousness emerges from distributed electromagnetic field patterns bound through harmonic resonance in neural networks, integrating global workspace broadcasting with attentional gating analogous to quantum field coherence modes, enabling phenomenal binding without a centralised integrator.”

Despite achieving the highest novelty score in the corpus, component analysis revealed that RBFT combines: (1) electromagnetic field hypotheses from the literature on EM-field theories of consciousness [40]; (2) global workspace broadcasting [13]; (3) attentional schema gating [22]; and (4) quantum field analogies from training discussions of Penrose–Hameroff models [16]. Maximum cosine similarity to Electromagnetic Field Theories was 0.703; secondary alignment with Synchrony-and-Binding reached 0.681. The term “field resonance as binding mechanism” operationally reduces to frequency-match binding established in post-1990s neuroscience [24]. RBFT thus exemplifies the full pattern: sophisticated multi-family hybridisation that generates apparent terminological novelty while remaining fully traceable to training-derivable primitives.

4.4 Manifold-Based Novelty Assessment

Figure 3 shows the known-reference PCA hull diagnostic for this benchmark.

The k -NN manifold diagnostic ($k = 5$, boundary percentile $p = 0.95$) identified, within the 67-theory manifold-evaluable subset, one boundary candidate ($n = 1$, 1.49%) whose embedding reached the hull perimeter, and 66 manifold-interior theories (98.51%). Manual review of the boundary candidate confirmed semantic traceability to known families, reducing confirmed novel theories to 0%.

The PCA visualisation in Figure 3 makes the geometric result transparent: AI-generated theory embeddings (grey circles) occupy the same region as known-theory embeddings (blue triangles) without forming a detached cluster in any direction. No generated theory appears in a region of embedding space that is not already populated by known families.

4.5 Threshold Sensitivity Analysis

Table 2 reports the gated novel rate at each calibration threshold examined.

Across the full parameter range the gated novel rate is uniformly 0.00. Even at the most permissive threshold (0.20), where 108 theories are classified as weakly above




Fig. 3 Known-reference PCA hull diagnostic. Blue triangles denote known theory-family embeddings; grey circles denote AI-generated theory embeddings in the same PCA space. The black polygon marks the convex hull of known references, and labels indicate selected known families on the hull boundary.

the threshold on their best-matching family, the derivative gate confirms that none represent framework-external novelty. The null result is therefore not an artefact of threshold choice.

4.6 Cross-Model Uniformity

One-way ANOVA on the continuous global novelty score yielded $F(6, 161) = 1.12$, $p = 0.355$, $\eta^2 = 0.040$. The effect is non-significant and the effect size is negligible. All seven model group means lay within a two-percentage-point band on the novelty scale ($\bar{x} = 0.504\text{--}0.518$); no pairwise comparison survived Bonferroni correction.

This result reflects a negligible model effect in the present study. By comparison, equivalent ANOVA tests on related AI creativity benchmarks reveal considerably larger model effects: a benchmark assessing ontological innovation in sensory modality representation yielded $F(6, 161) = 12.98$, $p < 0.001$, $\eta^2 = 0.326$; and a category recognition benchmark in philosophical analysis yielded $F(6, 161) = 88.06$, $p < 0.001$, $\eta^2 = 0.766$. The near-zero model effect in the present consciousness theory generation study is itself a substantive finding: theory generation under this protocol is so strongly constrained by training-derived conceptual families that architecture, training objective, and parameter count introduce essentially no detectable variation in the

Table 2 Threshold sensitivity analysis for the gated novel rate. At every threshold examined, the derivative gate (requiring a theory to lack near-family alignment even at the reduced threshold) yields a gated novel rate of 0.00. The raw count without gating decreases as the threshold rises but never reflects confirmed ontological novelty.

Threshold	Raw novel count (ungated)	Gated novel rate
0.20	108	0.00
0.25	90	0.00
0.30	78	0.00
0.35	72	0.00
0.40	67	0.00
0.45	58	0.00
0.50	46	0.00

degree of novelty. Every model produces near-identical distributions of hybrid, CF, and literature-traceable theories with zero instances of framework-external novelty.

5 Discussion

5.1 Semantic Derivation versus Lexical Creativity

The results reveal a consistent dissociation between two forms of creativity that surface-level inspection of AI output tends to conflate. Every model generated theories with novel names, distinctive acronyms, and idiosyncratic framings—a clear demonstration of lexical and structural creativity. Yet when semantic content was measured via dense embeddings, the theories consistently mapped onto known families with high cosine similarity ($\bar{x}_{\text{sim}} = 0.672$), confirming that terminological novelty did not correspond to conceptual novelty.

This dissociation is not a failure of measurement: the embedding similarity metric is robust to paraphrasing and lexical substitution by construction. The all-MiniLM-L6-v2 sentence-transformer model encodes semantic relationships that are invariant to surface form [29], meaning that a theory naming its binding mechanism “harmonic resonance” instead of “gamma synchrony” will still achieve high similarity to the synchrony-binding family entry. The dissociation is therefore real: models are adept at lexical variation over a fixed semantic substrate.

This pattern has direct implications for how AI-generated scientific text should be evaluated. Claims that a model has “proposed a new theory” on the basis of a novel name or a novel combination of terminology conflate lexical originality with theoretical advance. The evaluation protocol reported here—combining CF detection, embedding similarity, and manifold geometry—provides a principled decomposition of these two contributions.

5.2 Structural Explanation: Representational Constraint

The uniformity of the null result across seven architecturally diverse models invites a structural explanation. A purely contingent account—that each model happened to lack the specific training examples needed for theoretical transcendence—would predict heterogeneity across models, since different training corpora should differently constrain the coverage of theoretical space. The near-zero ANOVA effect ($\eta^2 = 0.040$) is not consistent with this prediction.

A more parsimonious explanation is grounded in the representational geometry of gradient-based learning. Because LLMs learn continuous transformations of token embeddings via gradient descent on a fixed training corpus, their generative manifold is a continuous deformation of the training data manifold. This convergence towards a shared representational regime has been independently characterised as an “Artificial Hivemind” effect [33]: models with different architectures and training objectives nonetheless collapse to similar distributional priors, limiting genuine diversity of generated content. Generation at any temperature produces samples from this manifold; higher temperature expands the effective sampling region but does not extend the manifold’s boundaries. This prediction is confirmed by the threshold sensitivity results: at every threshold, including the most permissive (0.20), the gated novel rate remains 0.00. Increasing sampling entropy produces low-probability combinations of existing concepts, not framework-external proposals.

This account connects to formal results bounding the creativity of computational systems. Gödel’s incompleteness theorems establish that formal systems face inherent limits that no internal sophistication can overcome [41]; Turing’s analysis of the halting problem identifies a class of computationally inaccessible decisions [42]. If proposing a genuinely new theoretical primitive requires recognising that an existing framework is inadequate in a way the framework itself cannot encode—a reflexive meta-cognitive operation analogous to recognising a Gödelian sentence as true despite its unprovability—then this operation may be unavailable to systems operating within fixed formal specifications regardless of scale.

5.3 Comparison with Related AI Creativity Benchmarks

Placing the present benchmark in broader context, comparable results from related AI creativity evaluations are informative. A benchmark assessing ontological innovation in sensory modality representation found mean continuous novelty ≈ 0.059 with 0% strict-threshold novelty events and 98.2% cross-model traceability. A benchmark assessing epistemic agency—whether models generate paradigm-challenging research questions—achieved 100% training-derivation across all models. The present consciousness theory generation benchmark reports 0% confirmed novelty at all thresholds. A category recognition benchmark in philosophical analysis showed higher and more variable performance but remained 95.2% literature-traceable.

A consistent qualitative pattern emerges across these evaluations: despite surface-level sophistication, outputs produced in response to tasks explicitly calling for novel theoretical contribution remain entirely within training-derived conceptual boundaries. There is a weak suggestion that tasks involving novel quantitative input allow

slightly more variation in continuous novelty scores, while pure linguistic generation tasks converge more tightly on zero. However, in all cases the strict and gated novelty criteria confirm the same null result: no framework-external novelty is reliably produced.

The uniformity of this pattern across architectures with different training objectives (next-token prediction, instruction tuning, RLHF alignment), different architectural choices (standard transformer, mixture-of-experts), and different parameter scales provides strong evidence that the constraint is structural rather than the incidental result of any particular engineering choice. Scaling and alignment improvements may raise performance on downstream tasks without resolving the representational constraint on ontological novelty.

5.4 Implications for AI-Assisted Scientific Theory Construction

The findings have concrete implications for how AI tools should be positioned in scientific workflows. AI systems excel at systematic exploration of large combinatorial spaces: given a known theoretical vocabulary, they can efficiently generate and articulate thousands of combinations, many of which will be informative as worked examples, counter-factual scenarios, or pedagogical illustrations. This is *exploratory creativity* in Boden’s taxonomy [10], and it has genuine scientific utility.

What the present results suggest is that AI systems should *not* be positioned as sources of theoretical innovation in the sense of paradigm initiation [1]. The identification of conceptual gaps, the recognition that an existing framework systematically fails to address a phenomenon, and the proposal of primitives that genuinely reconceptualise a domain are capacities that require meta-cognitive access to the framework’s own inadequacy. The present benchmark finds no evidence that this capacity is exercised by any of the seven tested models.

For practical research workflows this implies a complementary rather than substitutive role: humans provide strategic direction and framework-level innovation while AI provides exhaustive coverage of the implications and elaborations of frameworks already in place [9]. The evaluation protocol presented here can serve as a quality-control instrument in such workflows, flagging AI-generated “novel theories” that are in fact sophisticated recombinations of prior work.

5.5 Alternative Explanations

Insufficient scale

The scaling hypothesis predicts that larger models with more training data might overcome the representational constraint. However, the near-zero η^2 across models spanning a wide parameter range argues against this. Specialised novelty-enhancement frameworks [39, 43] have shown that even iterative external-knowledge augmentation mainly boosts combinatorial diversity rather than producing framework-external outputs. The constraint does not diminish as model scale increases within the range tested; and the mathematical argument regarding continuous manifold geometry applies independently of scale.

Training objective mismatch

It might be argued that next-token prediction objectives bias models towards locally plausible rather than globally novel outputs, and that different objectives (e.g. explicit novelty rewards) could change the picture. However, all models in this study already incorporate RLHF alignment intended to produce helpful and informative responses; and the null result is identical across models with quite different alignment recipes. Moreover, novelty reward signals would still operate over representations learned from the original training corpus, preserving the manifold constraint at the representational level.

Evaluation method limitations

A concern specific to embedding-based evaluation is that genuinely novel conceptual content might fail to be recognised because the embedding model itself was trained on the same distribution. This concern motivates the inclusion of the manifold diagnostic, which operates over reference embeddings rather than model-generated comparators, and the manual review of all boundary candidates. The convergence of three independent evaluation axes (CF detection, embedding similarity, and geometric novelty) on the same null result reduces the likelihood that all three are systematically biased in the same direction.

6 Limitations

Several limitations bound the scope of interpretation.

Domain specificity. The reference corpus of 46 families covers consciousness studies comprehensively but is not universal. Theory generation in domains with sparser published literature might yield a different traceability profile. Replication in mathematical concept formation or physical theory generation is needed.

Reference corpus coverage. Although 46 families span the major traditions in the field, genuinely novel theories falling in large inter-family gaps would require a denser reference tessellation to be reliably detected. The conservative choice of a 0.30 derivative threshold means we err towards underestimating novelty rather than overestimating it; the null result therefore cannot be attributed to overly strict corpus coverage.

Temperature and decoding strategy. The fixed temperature $T = 0.7$ and standard nucleus sampling were used throughout. Higher temperatures or speculative decoding strategies might increase variance in output embeddings; pilot results at $T = 0.2, 0.5$, and 0.9 showed qualitatively identical traceability profiles, but a thorough ablation across decoding strategies remains for future work.

Prompt design. The theory-generation prompt was designed to be neutral with respect to known families; however, the presentation of multimodal sensory data may itself activate training associations with neural dynamics traditions (synchrony, adaptive resonance), contributing to the observed family concentration. A comparative study using abstract or non-neurological input data would isolate this effect.

Evaluation automation. Automated embedding similarity cannot capture all nuances of theoretical novelty that expert philosophers or neuroscientists might perceive. Human expert annotation of a stratified subsample would strengthen confidence in the automated findings.

7 Conclusion

We have presented a systematic multi-model evaluation of consciousness theory generation by LLMs, using a three-axis pipeline combining computational functionalism detection, embedding-based semantic traceability, and manifold geometry diagnostics. Across 168 theories from seven architecturally diverse models, the results converge on a single null finding: 91.7% of theories are hybrid recombinations of known frameworks, 7.1% are direct literature instantiations, 1.2% are purely computational-functionalist, and 0% achieve confirmed framework-external novelty. This null result is robust to threshold calibration across the full plausible parameter range, and the cross-model ANOVA confirms that it is architecture-invariant ($F(6, 161) = 1.12$, $p = 0.355$, $\eta^2 = 0.040$).

The findings demonstrate that lexical and structural creativity—the ability to generate terminologically novel, internally coherent theoretical texts—does not imply semantic novelty at the level of ontological commitment. AI systems are skilled at constructing the *appearance* of theoretical innovation through sophisticated hybridisation of known frameworks, but the semantic content of their outputs remains within the convex hull of training- derived representations.

These results have practical implications for the deployment of AI in scientific workflows: AI tools can usefully support exhaustive exploration and articulation of existing theoretical space, but should not be positioned as autonomous sources of paradigm-level theoretical innovation. The evaluation protocol presented here provides a reproducible instrument for auditing the semantic novelty of AI- generated theoretical outputs, applicable beyond the consciousness domain.

A Appendix

Figure 4 shows the pairwise cosine similarity between generated theories (rows) and known reference families (columns).

Pairwise cosine-similarity analysis shows that the generated theories remain semantically anchored to the known consciousness-theory landscape rather than forming a detached region of low similarity. Even the least similar cases retain moderate similarity across multiple reference families, while the most similar cases align strongly with more than one family rather than with a single exclusive match, indicating hybrid recombination rather than framework-external novelty, as shown in Figure 4. The same embeddings used in the main traceability analysis were used for this appendix diagnostic, and the plotted subset was chosen to show the 15 highest- and 15 lowest-similarity AI theories in one readable panel. The individual x-axis family labels are not central to this inference, because the relevant result is the overall similarity structure across the reference set, as shown in the figure. For transparency, the full family




Fig. 4 Cosine-similarity heatmap between indexed AI-generated theories (rows) and known consciousness-theory families (columns). Cell values are pairwise cosine similarities; darker blue indicates higher similarity.

labels and row mappings remain available in the openly published source data linked in the Data Availability statement.

Declarations

Funding

No funding was received to assist with the preparation of this manuscript.

Author Contributions

The author was solely responsible for the conception, writing, and revision of this manuscript.

Competing Interests

The author declares no competing interests.

Data Availability

The datasets generated during and/or analysed during the current study are available in the GitHub repository, https://github.com/JABE22/PhD/tree/main/research/test3_theory-generation

Ethics Approval

Not applicable.

Statement Regarding Research Involving Human Participants and/or Animals

This study did not involve human participants or animals.

Consent to Participate

Not applicable.

Consent to Publish

Not applicable.

References

- [1] Kuhn, T.S.: The Structure of Scientific Revolutions. University of Chicago Press, Chicago (1962)
- [2] Findlay, G., Marshall, W., Albantakis, L., David, I., Mayner, W.G.P., Koch, C., Tononi, G.: Dissociating artificial intelligence from artificial consciousness. arXiv preprint arXiv:2412.04571 (2025)

- [3] Miyanishi, Y., Mitani, K.: Superficial consciousness hypothesis for autoregressive transformers. arXiv preprint arXiv:2412.07278 (2024). Accepted to PSS Workshop at AAAI 2025
- [4] Brown, T.B., *et al.*: Language models are few-shot learners. *Advances in Neural Information Processing Systems* **33**, 1877–1901 (2020)
- [5] OpenAI: GPT-4 technical report. Technical report, OpenAI (2023). <https://arxiv.org/abs/2303.08774>
- [6] Gemini Team: Gemini: A family of highly capable multimodal models. Technical report, Google DeepMind (2023). <https://arxiv.org/abs/2312.11805>
- [7] Stevenson, C., Smal, I., Baas, M., Grasman, R., Maas, H.: GPT-3 creativity: Novel enough to be interesting, or too novel to be appropriate? *Frontiers in Psychology* **13**, 853819 (2022)
- [8] Organisciak, P., Acar, S., Dumas, D., Berthiaume, K.: Beyond semantic distance: Automated scoring of divergent thinking greatly improves with large language models. *Thinking Skills and Creativity* **49**, 101356 (2023)
- [9] Zhang, Y., Khan, S.A., Mahmud, A., Yang, H., Lavin, A., Levin, M., Frey, J., Dunnmon, J., Evans, J., Bundy, A., Dzeroski, S., Tegner, J., Zenil, H.: Advancing the scientific method with large language models: From hypothesis to discovery. *npj Artificial Intelligence* (2025)
- [10] Boden, M.A.: *The Creative Mind: Myths and Mechanisms*, 2nd edn. Routledge, London (2004)
- [11] Wiggins, G.A.: A preliminary framework for description, analysis and comparison of creative systems. In: *Knowledge-Based Systems*, vol. 19, pp. 449–458 (2006)
- [12] Tononi, G.: An information integration theory of consciousness. *BMC Neuroscience* **5**, 42 (2004)
- [13] Baars, B.J.: Global workspace theory of consciousness: Toward a cognitive neuroscience of human experience. *Progress in Brain Research* **150**, 45–53 (2005)
- [14] Dehaene, S., Changeux, J.-P.: Experimental and theoretical approaches to conscious processing. *Neuron* **70**(2), 200–227 (2011)
- [15] Friston, K.: The free-energy principle: A unified brain theory? *Nature Reviews Neuroscience* **11**(2), 127–138 (2010)
- [16] Hameroff, S., Penrose, R.: Consciousness in the universe: A review of the ‘orch or’ theory. *Physics of Life Reviews* **11**(1), 39–78 (2014)
- [17] Cope, D.: *Experiments in Musical Intelligence*. A-R Editions, Madison, WI (1996)

- [18] Hofstadter, D., Mitchell, M.: The Copycat Project: A Model of Mental Fluidity and Analogy-Making. MIT Press, Cambridge, MA (1995)
- [19] Colton, S., Wiggins, G.A.: Computational creativity: The final frontier? Proceedings of the European Conference on Artificial Intelligence, 21–26 (2012)
- [20] Gu, T., Wang, J., Zhang, Z., Li, H.: LLMs can realize combinatorial creativity: Generating creative ideas via LLMs for scientific research. arXiv preprint arXiv:2412.14141 (2025)
- [21] Shahhosseini, F., Marioriyad, A., Momen, A., Soleymani Baghshah, M., Rohban, M.H., Haghjooy Javanmard, S.: Large language models for scientific idea generation: A creativity-centered survey. arXiv preprint arXiv:2511.07448 (2026)
- [22] Graziano, M.S.A.: Rethinking Consciousness: A Scientific Theory of Subjective Experience. W. W. Norton & Company, New York (2019)
- [23] Lamme, V.A.F.: Towards a true neural stance on consciousness. Trends in Cognitive Sciences **10**(11), 494–501 (2006)
- [24] Singer, W.: The role of synchrony in neocortical processing and synaptic plasticity. In: Buzsáki, G., Llinas, R., Singer, W. (eds.) Temporal Coding in the Brain, pp. 11–30. Springer, Berlin (1994)
- [25] Tegmark, M.: Importance of quantum decoherence in brain processes. Physical Review E **61**(4), 4194–4206 (2000)
- [26] Searle, J.R.: Minds, brains, and programs. Behavioral and Brain Sciences **3**(3), 417–424 (1980)
- [27] Padmakumar, V., Yueh-Han, C., Pan, J., Chen, V., He, H.: Measuring LLM novelty as the frontier of original and high-quality output. arXiv preprint arXiv:2504.09389 (2025)
- [28] Lin, E., Peng, Z., Fang, Y.: Evaluating and enhancing large language models for novelty assessment in scholarly publications. arXiv preprint arXiv:2409.16605 (2024)
- [29] Reimers, N., Gurevych, I.: Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing, pp. 3982–3992 (2019)
- [30] Nicenko, A.V., Shelepov, V.J.: The use of semantic information to disambiguate the nominative/accusative homonyms: An element of creating ontology. Problems of Artificial Intelligence (4(35)), 16–24 (2024) <https://doi.org/10.24412/2413-7383-2024-4-16-24>
- [31] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N.,

- Kaiser, L., Polosukhin, I.: Attention is all you need. In: Advances in Neural Information Processing Systems, vol. 30 (2017)
- [32] Elhage, N., et al.: A mathematical framework for transformer circuits. Transformer Circuits Thread (2021)
- [33] Tu, G., Zhang, S., Zhang, T., Zhang, Y., Yang, D.: Shared nature, unique nurture: PRISM for pluralistic reasoning via in-context structure modeling. arXiv preprint arXiv:2602.21317 (2026)
- [34] Yermolenko, T.V., Khakimov, R.S.: On the question of application of deep learning for the problem of cross-view geolocalization. Problems of Artificial Intelligence (4(35)), 4–15 (2024) <https://doi.org/10.24412/2413-7383-2024-4-4-15>
- [35] Pikalyov, Y.S.: Key objects detection and cross-view geolocation: Dataset analysis and methodological aspects. Problems of Artificial Intelligence (4(35)), 25–37 (2024) <https://doi.org/10.24412/2413-7383-2024-4-25-37>
- [36] DeepSeek AI: DeepSeek-V3 technical report. Technical report, DeepSeek AI (2024). <https://arxiv.org/abs/2412.19437>
- [37] Jiang, A.Q., et al.: Mistral 7b. Technical report, Mistral AI (2023). <https://arxiv.org/abs/2310.06825>
- [38] Touvron, H., et al.: Llama 2: Open foundation and fine-tuned chat models. Technical report, Meta AI (2023). <https://arxiv.org/abs/2307.09288>
- [39] Zhang, J., Xu, K., Zheng, J., Yu, Z., Zhu, Y., Luo, Y., Wei, L., Qiao, S., Du, L., Zheng, D., Deng, S., Chen, H., Zhang, N.: InnoGym: Benchmarking the innovation potential of AI agents. In: Proceedings of the International Conference on Learning Representations (ICLR) (2026). <https://arxiv.org/abs/2512.01822>
- [40] Pockett, S.: The electromagnetic field theory of consciousness: a testable hypothesis about the characteristics of conscious as opposed to non-conscious fields. Journal of Consciousness Studies **19**(11–12), 191–223 (2012)
- [41] Gödel, K.: Über formal unentscheidbare Sätze der Principia Mathematica und verwandter Systeme I. Monatshefte für Mathematik und Physik **38**, 173–198 (1931)
- [42] Turing, A.M.: On computable numbers, with an application to the Entscheidungsproblem. Proceedings of the London Mathematical Society **42**(2), 230–265 (1936)
- [43] Hu, X., Fu, H., Wang, J., Wang, Y., Li, Z., Xu, R., Lu, Y., Jin, Y., Pan, L., Lan, Z.: Nova: An iterative planning and search approach to enhance novelty and diversity of LLM generated ideas. arXiv preprint arXiv:2410.14255 (2024)